

Dell PowerStore: Best Practices Guide

May 2024

H18241.7

White Paper

Abstract

This document provides best practices for installing and configuring Dell PowerStore for optimal performance and availability.

Copyright

The information in this publication is provided as is. Dell Inc. makes no representations or warranties of any kind with respect to the information in this publication, and specifically disclaims implied warranties of merchantability or fitness for a particular purpose.

Use, copying, and distribution of any software described in this publication requires an applicable software license.

Copyright © 2020-2024 Dell Inc. or its subsidiaries. Published in the USA May 2024 H18241.7.

Dell Inc. believes the information in this document is accurate as of its publication date. The information is subject to change without notice.

Contents

Executive summary.....4

Introduction4

Hardware considerations.....7

Network considerations.....9

PowerStore storage resources.....12

PowerStore features and layered applications14

External host considerations16

Conclusion.....17

References.....18

Executive summary

Overview

This white paper provides best practices guidance for using Dell PowerStore in a mixed-business environment. It focuses on optimizing system performance and availability, and on maximizing the usability of the automated storage features.

These guidelines are intended to cover most use cases. They are recommended by Dell Technologies but are not strictly required. Some exception cases are addressed in this white paper. Less common edge cases are not covered by these general guidelines and are addressed in use case-specific white papers.

For questions about the applicability of these guidelines in a specific environment, contact your Dell Technologies representative to discuss the recommendations.

Audience

This document is intended for IT administrators, storage architects, partners, and Dell Technologies employees. This audience also includes any individuals who may evaluate, acquire, manage, operate, or design a Dell networked storage environment using PowerStore systems.

Revisions

Date	Part number/ revision	Description
April 2020		Initial release: PowerStoreOS 1.0
June 2020		Terminology updates
September 2020		Terminology updates
June 2021		PowerStoreOS 2.0 updates
January 2022		PowerStoreOS 2.1 updates; template update
July 2022	H18241.5	PowerStoreOS 3.0 updates
May 2023	H18241.6	PowerStoreOS 3.5 updates
May 2024	H18241.7	PowerStoreOS 4.0 updates Removed references to PowerStore X

Introduction

Document purpose

This document introduces specific configuration recommendations that enable optimal performance from PowerStore.

PowerStore overview

PowerStore achieves new levels of operational simplicity and agility. It uses a container-based microservices architecture, advanced storage technologies, and integrated machine learning to unlock the power of your data. PowerStore is a versatile platform with a performance-centric design that delivers multidimensional scale, always-on data reduction, and support for next-generation media.

PowerStore brings the simplicity of public cloud to on-premises infrastructure, streamlining operations with an integrated machine-learning engine and seamless automation. It also offers predictive analytics to easily monitor, analyze, and troubleshoot the environment. PowerStore is highly adaptable, providing the flexibility to host specialized workloads directly on the appliance and to modernize infrastructure without disruption. It also offers investment protection through flexible payment solutions and data-in-place upgrades.

Terminology

The following table provides definitions for some of the terms that are used in this document.

Table 1. Terminology

Term	Definition
Appliance	A solution containing a base enclosure and any attached expansion shelves. The size of an appliance could be the base enclosure only or the base enclosure plus expansion shelves.
Base enclosure	Refers to the enclosure that contains both nodes (node A and node B) and 25 NVMe drive slots.
Cluster	One or more appliances in a single grouping and management interface. Clusters are expandable by adding more appliances to the existing cluster, up to the allowed amount for a cluster.
Embedded module	Connectivity card in the PowerStore node that provides ports for Ethernet connections, and various service and management ports.
Expansion enclosure	Enclosures that can be attached to a base enclosure to provide additional storage in the form of SAS drives.
Fibre Channel	A protocol used to perform NVMe or SCSI commands over a Fibre Channel (FC) network.
File system	A storage resource that can be accessed through file-sharing protocols such as SMB or NFS.
Internet SCSI (iSCSI)	Provides a mechanism for accessing block-level data storage over network connections.
I/O module	Optional connectivity cards that provide additional Fibre Channel or Ethernet ports.
IOPS	I/Os per second, a measure of transactional performance for small-block workloads.
MBPS	Megabytes per second, a measure of bandwidth performance for large-block workloads.
Network-attached storage (NAS) server	A virtualized network-attached storage server that uses the SMB, NFS, or FTP/SFTP protocols to catalog, organize, and transfer files within file system shares and exports. A NAS server, the basis for multi-tenancy, must be created before creating file-level storage resources. A NAS server is responsible for the configuration parameters on the set of file systems that it serves.
Network File System (NFS)	An access protocol that allows data access from Linux/UNIX hosts on a network.

Term	Definition
Node	A storage node that provides the processing resources for performing storage operations and servicing I/O between storage and hosts.
NVMe over Fibre Channel (NVMe/FC)	Allows hosts to access storage systems across a Fibre Channel network fabric using the NVMe protocol.
NVMe over TCP (NVMe/TCP)	Allows hosts to access storage systems across a TCP network fabric using the NVMe protocol.
PowerStore Command Line Interface (PSTCLI)	An interface that allows a user to perform tasks on the storage system by typing commands instead of using the user interface (UI).
PowerStore Q model	Container-based storage system that is running on purpose-built hardware. This storage system supports unified (block and file) workloads, or block-optimized workloads. The PowerStore Q model supports Quad-Level Cell (QLC) NVMe SSDs for data storage.
PowerStore T model	Container-based storage system that is running on purpose-built hardware. This storage system supports unified (block and file) workloads or block optimized workloads. The PowerStore T model supports Triple-Level Cell (TLC) NVMe SSDs for data storage.
Server Message Block (SMB)	An access protocol that allows remote file data access from clients to hosts on a network. This protocol is typically used in Microsoft Windows environments.
Snapshot	A point-in-time view of data stored on a storage resource. A user can recover files from a snapshot, restore a storage resource from a snapshot, or provide access to a host.
Thin clone	A read/write copy of a volume, volume group, file system, NAS server, or snapshot that shares blocks with the parent resource.
Virtual Volumes (vVols)	A VMware storage framework which allows VM data to be stored on individual VMware vSphere Virtual Volumes (vVols). This framework allows for data services to be applied at a VM-granularity level while using Storage Policy Based Management (SPBM).
Volume	A block-level storage device that can be shared out using a protocol such as iSCSI or Fibre Channel.

We value your feedback

Dell Technologies and the authors of this document welcome your feedback on this document. Contact the Dell Technologies team by [email](#).

Author: Ethan Stokes

Contributors: Stephen Wright, Andrew Sirpis, Ryan Poulin

Note: For links to other documentation for this topic, see the [PowerStore Info Hub](#).

Hardware considerations

Introduction

At the highest level, design for optimal performance follows a few simple rules. The main principles of designing a PowerStore system for performance are:

- Distribute workloads across available resources
- Simplify the configuration
- Design for resilience
- Maintain the latest-released PowerStoreOS version

Hardware components are the foundation of any storage system. This section discusses some key hardware differences between PowerStore models that help determine performance. It also explains how different configuration options can result in different performance from the same hardware.

PowerStore deployment modes

The PowerStore platform consists of eleven different models, from the PowerStore 500 model through the PowerStore 9200 model. PowerStoreOS 1.0 launched with the PowerStore 1000T, 3000T, 5000T, 7000T, and 9000T. PowerStoreOS 2.0 introduced the entry-level 500. PowerStoreOS 3.0 introduced a platform refresh, with the 1200T, 3200T, 5200T, and 9200T models. PowerStoreOS 4.0 introduced the 3200Q model appliance. All models use a common base enclosure and I/O modules. The models differ by CPU core count and speed, memory size, and number of NVMe NVRAM drives. These hardware differences give each model a unique performance profile.

Besides the hardware differences between the models, PowerStore can be installed in one of two different deployment modes. Each deployment mode has different capabilities, as detailed in Table 2. Choose the deployment mode that provides the required capabilities.

Table 2. PowerStore configurations

Deployment mode	External block access	External file access
Unified	✓	✓
Block optimized	✓	✗

The PowerStore system has different performance characteristics depending on deployment mode.

PowerStore Deployment Modes

PowerStore appliances can be installed in a unified configuration that provides file and block access, or in a block optimized configuration that provides only block access.

Deployment Mode : Unified

PowerStore appliances with unified modes can provide access to block and file storage resources simultaneously. This is the default deployment mode.

Deployment Mode : Block optimized

If the PowerStore appliance will not be used for file access, it can be installed in block optimized mode, which disables the file capabilities. This mode can increase the amount of block workload that the system can provide, because it can devote the additional CPU and memory that is no longer needed for file capabilities.

Relative performance expectations

This section compares the performance potential of the different PowerStore models when serving external workload. Performance scales based on the specific hardware complement of the model and is also impacted by the configuration type.

In general, the IOPS capability of the PowerStore models scales linearly from PowerStore 500 up to 9200 models. As mentioned previously, deployment mode also impacts performance capability. A PowerStore T model in block optimized mode can deliver more block IOPS than the same model in unified mode.

Except for PowerStore 500, PowerStore systems use NVMe NVRAM drives to provide persistent storage for cached write data. PowerStore 1000 up to 3200 model arrays have two NVRAM drives per system, while PowerStore 5000 up to 9200 model arrays have four NVRAM drives per system. The extra drives mean that these systems can provide higher MBPS for large-block write workloads.

PowerStore cluster

PowerStore systems can be clustered. A PowerStore cluster combines multiple PowerStore appliances into a single grouping that is managed as a single storage system. A PowerStore cluster delivers aggregate performance from all appliances in the cluster, but a single volume is serviced by only one appliance at any given time. While not required, it is recommended that all appliances in a cluster be of the same model and have similar physical capacities to provide consistent performance across the cluster. PowerStore T models and PowerStore Q models can be mixed within the same cluster.

Volumes can be migrated between appliances in a cluster. It is recommended that any host that is connected to a PowerStore cluster has equivalent connectivity to all appliances in the cluster. All appliances in a cluster should be physically located in the same data center and must be connected to the same LAN.

Clustering is applicable to block storage resources only. While a PowerStore appliance in Unified mode can serve as the cluster's primary appliance, the file resources cannot migrate to a different appliance. When deploying multiple appliances for file access, plan to have multiple clusters.

Drive configuration

PowerStore can be configured with NVMe solid-state devices (SSDs) or NVMe storage class memory (SCM) drives for user data. The 3200Q model only supports QLC NVMe SSDs, while all other PowerStore model appliances support TLC NVMe SSDs. Mixing QLC and TLC drives within a single appliance is not supported. SSD-based systems can be expanded with additional drives to increase the amount of available storage capacity. PowerStore 1000, 3000, 5000, 7000, and 9000 models can be expanded with SAS SSD-based expansion shelves. Starting in PowerStoreOS 3.0, all PowerStore models, including the 3200Q, can be expanded with NVMe SSD-based expansion shelves if they meet the necessary hardware prerequisites defined in the white paper [Dell PowerStore:](#)

[Introduction to the Platform](#). It is recommended that all drives within a PowerStore system be the same size, which can maximize the usable capacity from each drive.

PowerStore Dynamic Resiliency Engine (DRE) is used to manage the drives in the system. All drives are automatically used to provide storage capacity. DRE groups the drives into resiliency sets to protect against drive failure. User configuration of the drives is not necessary, and dedicated hot spare drives are not required in PowerStore. Spare space for rebuilds is automatically distributed across all drives within each resiliency set. This configuration provides better resource utilization and enables faster rebuilds if there is a drive failure.

At initial installation of the PowerStore system, DRE can be configured with either single- or double-drive failure tolerance. To provide the greatest usable capacity from the same number of drives, it is recommended to initially install PowerStore with a minimum of ten drives for single-drive failure tolerance, or nineteen drives for double-drive failure tolerance. PowerStore 3200Q models require a minimum of eleven drives, and double-drive failure tolerance is enforced on the appliance.

SCM drives

PowerStore can use SCM drives either by having only SCM drives installed in the system, or, with PowerStoreOS 2.0 or higher, by mixing SCM and TLC model SSD drives.

Systems with all SCM drives are recommended for small-block workloads that require the absolute lowest latencies. A system with all SCM drives will place both data and metadata on the SCM drives.

Systems with mixed SSD and SCM drives will use the SCM drives for metadata acceleration; the SCM drives will store metadata for faster lookups. This can reduce latency on read operations in systems with large physical capacities. When mixing SCM and SSDs, it is recommended that at least 5 percent of the system's physical capacity should be SCM.

Network considerations

Introduction

External hosts send and receive data from PowerStore through Fibre Channel, Ethernet, or both networks. These networks play a large role in determining the performance potential of PowerStore. This section discusses considerations for the external network, and for the PowerStore network ports.

Host networking recommendations are covered in [External host considerations](#).

General network performance and high availability

It is recommended to use redundant switch hardware between the PowerStore system and external clients. For more details about configuring a redundant network, see the document *PowerStore Host Configuration Guide* on [Dell.com/powerstoredocs](https://dell.com/powerstoredocs).

Fibre Channel fabrics

For Fibre Channel connectivity, configure dual redundant fabrics, with each PowerStore node and each external host having connectivity on each of the fabrics. Minimize the number of hops between host and PowerStore.

For performance, load balancing, and redundancy, each host should have at least two paths to each PowerStore node (four paths per PowerStore appliance). It is recommended that a host should have no more than eight paths per volume.

Ethernet networks

For Ethernet connectivity, use multiple switches that are connected with Virtual Link Trunking interconnect (VLTi) and Link Aggregation Control Protocol (LACP) or equivalent technologies. Each PowerStore node should have connectivity to all linked switches.

The first two ports of the embedded module 4-port card on each PowerStore node are bonded together within the PowerStoreOS. For highest performance and availability from these ports, it is recommended also to configure link aggregation across the corresponding switch ports.

With the introduction of PowerStoreOS 3.0, user-defined link aggregations are supported for Ethernet ports. Link aggregations can include 2 to 4 different Ethernet ports. The ports must be on the same node and operate at the same speed. A mirror link aggregation will automatically be created on the peer node. In PowerStoreOS releases prior to 4.0, user-defined link aggregations only support NAS server interfaces. In PowerStoreOS 4.0 and later, user-defined link aggregations also support iSCSI host connectivity and replication.

When using Ethernet for block access to PowerStore volumes (by means of iSCSI and/or NVMe over TCP), each host should have at least two paths to each PowerStore node (four paths per PowerStore appliance). It is recommended that a host should have no more than eight paths per volume.

PowerStore front-end ports

PowerStore supports Ethernet connectivity through ports on the embedded module, and on optional I/O modules. PowerStore supports Fibre Channel connectivity through ports on optional I/O modules.

On PowerStore 1000 – PowerStore 9200 models, the I/O module in slot 0 is 16-lane PCIe Gen3, while the I/O module in slot 1 is 8-lane. For this reason, if the system is configured with a 100 GbE I/O Module, it must be installed in slot 0. When a Fibre Channel I/O module is being installed, it is recommended to always use I/O module slot 0 first unless the system contains a 100 GbE I/O Module. Both I/O modules on PowerStore 500 models are 8-lane PCIe, and the 100 GbE I/O module is not supported on this platform.

PowerStore Fibre Channel ports

PowerStore Fibre Channel ports support speeds for 32 Gb/s, 16 Gb/s, 8 Gb/s, and 4 Gb/s. This speed depends on the SFP used and the switchport or HBA that is connected. Because higher speeds allow for greater MBPS and IOPS capabilities, it is recommended that you use the highest speed supported by the environment.

Fibre Channel ports are available on I/O modules that are inserted into I/O module slots on the nodes. The Fibre Channel I/O module is 16-lane PCIe Gen3. On PowerStore 1000 – 9200 models, I/O module slot 0 is also 16-lane, while I/O module slot 1 is 8-lane. If Fibre Channel I/O modules are installed in both I/O module slots, it is recommended to cable the ports in I/O module slot 0 first, due to the PCIe difference. The PCIe lanes in I/O module slot 1 are only a limiting factor for total MBPS, and only when all four ports on the Fibre Channel I/O module are operating at 32 Gb/s. Both I/O module slots on PowerStore 500 are 8-lane PCIe, and therefore there is no slot preference.

NVMe over Fibre Channel

The NVMe over Fibre Channel (NVMe/FC) protocol provides connectivity using the same Fibre Channel ports, but can decrease the transport latency between PowerStore and the host. Note that all parts of the network, including switches and HBAs, must support NVMe over Fibre Channel.

PowerStore Ethernet ports

PowerStore optical Ethernet ports support speeds of up to 25 Gb/s, based on the SFP that is used. Copper Ethernet ports support speeds of up to 10 Gb/s. Because higher speeds allow for greater MBPS and IOPS capabilities, it is recommended that you use the highest speed supported by your environment.

With PowerStoreOS 3.0, a new 2-Port Ethernet card is introduced that supports speeds of up to 100 Gb/s. This 100 GbE card is supported on PowerStore 1000-9200 models in I/O Module 0 slot.

Jumbo frames (MTU 9000) are recommended for increased network efficiency. Jumbo frames must be supported on all parts of the network between PowerStore and the host.

Ethernet ports for iSCSI

Map additional Ethernet ports for iSCSI to increase system MBPS capabilities. Enable Jumbo frames for iSCSI by setting the Cluster MTU to 9000 and set the storage network MTU to 9000.

The embedded module 4-port card and the optional network I/O modules are 8-lane PCIe Gen3. When more than two 25 GbE ports are used, these cards are oversubscribed for MBPS. To maximize MBPS scaling in the system, consider cabling and mapping the first two ports of all cards in the system first. Then, cable and map other ports as needed.

When PowerStore models that are in unified mode are used for both iSCSI and file access, it is recommended that you use different physical ports for both NAS and iSCSI.

Ethernet ports for NVMe over TCP

The NVMe over TCP (NVMe/TCP) protocol provides connectivity using the same physical Ethernet ports as iSCSI. NVMe/TCP can be enabled on the same Storage Network as iSCSI, or different Storage Networks can be created to isolate iSCSI and NVMe/TCP traffic.

The embedded module 4-port card and the optional network I/O modules are 8-lane PCIe Gen3. When more than two 25 GbE ports are used, these cards are oversubscribed for MBPS. To maximize MBPS scaling in the system, consider cabling and mapping the first two ports of all cards in the system first. Then, cable and map other ports as needed.

When PowerStore models that are in unified mode are used for both NVMe/TCP and file access, it is recommended that you use different dedicated physical ports for both NAS and NVMe/TCP.

SmartFabric Storage Software (SFSS)

Dell SmartFabric Storage Software (SFSS) provides Centralized Discovery Controllers (CDCs) for NVMe/TCP Endpoints. These CDCs facilitate endpoint discovery, registration, soft zoning, and event notifications. With SFSS, Dell Technologies provides the industry's first comprehensive connectivity automation solution for NVMe/TCP endpoints such as

Dell PowerEdge, Dell PowerStore, and Dell PowerMax. For more information about SFSS support, see the [SmartFabric Storage Software Deployment Guide](#) on the [Storage Networking Info Hub](#).

Ethernet ports for NAS

It is recommended that you use bonded ports for NAS connectivity. Prior to PowerStoreOS 3.0, NAS servers automatically created their interfaces on the two bonded ports on the embedded module 4-port card. With PowerStoreOS 3.0, user-defined link aggregations can be used to reserve different physical ports for file access only. In PowerStoreOS 4.0, user-defined link aggregations also support storage iSCSI and replication connectivity. For highest performance and availability from any aggregated ports, it is recommended that you configure link aggregation across the corresponding switch ports.

Enable Jumbo frames for NAS by setting the cluster MTU to 9000.

If the PowerStore is also providing block access through iSCSI or NVMe-TCP, or asynchronous replication over Ethernet, it is recommended that you use different physical ports for NAS than the ports which are tagged for replication or storage networks.

Fail-Safe Networking (FSN)

PowerStoreOS 3.5 adds Fail-Safe Networking (FSN) support for file interfaces. FSN is a high-availability feature that enables configuring ports in a primary/backup configuration. Under normal circumstances, the primary ports are designated as active and are used to service IO. If all primary ports of an FSN go offline, the backup ports automatically become active and continue to service IO. This enables redundancy in case of port, cable, or switch failure. When the primary ports are restored, the system automatically makes the primary ports active again. For optimal performance in the event of a failure, it is suggested that the configurations of the active and standby ports or bonds selected be consistent. For more information about FSN, see the [Dell PowerStore: File Capabilities](#) white paper on the [PowerStore Info Hub](#).

PowerStore storage resources

Introduction

PowerStore allows access to block, file, or both storage types. This section provides recommendations for the different types of storage resources.

Block storage resources

Block storage resources are accessed through Fibre Channel, NVMe over Fibre Channel, iSCSI, or NVMe over TCP protocols. A host should only access a block resource using one of these protocols at a time. It is not supported for the same host to access the same block resource using more than one protocol.

Appliance balance for block workloads

PowerStore block storage resources are accessed using ALUA/ANA active/optimized or active/non-optimized paths between the host and the two nodes within the PowerStore appliance. I/O is normally sent on an active/optimized path. PowerStore automatically chooses one of the nodes for the active/optimized path, when the volume is mapped to the host, to maintain a balanced workload across the nodes. This characteristic is called **node affinity** and can be viewed in PowerStore Manager and modified with PSTCLI or

REST. These changes take effect immediately and are nondisruptive if the host is correctly configured for multipathing.

Dynamic Node Affinity

Introduced in PowerStoreOS 2.1, the node affinity of block storage resources will be dynamically rebalanced between nodes to maintain relatively consistent utilization, latency, and performance between both nodes of an appliance.

This feature works on block resources that have not had the node affinity manually set by means of PSTCLI or REST. If the node affinity was manually set, the volume must be unmapped and then remapped to the host, which will reset the affinity back to “system selected”. This impacts multipathing only; the system does not need to trespass any volume between nodes.

In PowerStoreOS 4.0, dynamic node affinity is expanded to include virtual volumes (vVols). Data and config vVols are eligible for rebalancing, while unbound vVols are ignored. When an imbalance is detected, the node affinity of volumes is changed first. If additional balancing is required after volumes are balanced, vVols are changed next.

This feature is designed to be transparent and seamless to clients. The goal is to achieve an even throughput, bandwidth, and CPU utilization balance across both nodes within an appliance, regardless of resource type. The system makes changes automatically without requiring any admin interaction. This provides balanced hardware utilization, optimizes performance, and adapts dynamically as workloads change.

Performance policy

All block storage resources in a PowerStore system have a defined performance policy. By default, this policy is set to Medium. The performance policy does not have any impact on system behavior unless some volumes have been set to Low Performance Policy, and other volumes are set to Medium or High. During times of system resource contention, PowerStore devotes fewer compute resources to volumes with Low Performance Policy. Reserve the Low policy for volumes that have less-critical performance needs.

Quality of Service (QoS)

In PowerStoreOS version 4.0, Quality of Service (QoS) can be configured for volumes and volume groups. A QoS policy, which includes an I/O limit rule, sets the maximum limits on I/O for volumes and volume groups. These policies are used to ensure that critical applications get priority over other workloads and provide predictable performance for each application. QoS policies are interoperable with existing performance policies.

As this feature is an I/O limiting feature, host performance can be impacted. The impact on host I/O depends directly on the amount of I/O to the resource and the limit configured. Having a limit too low can cause performance related issues, which can be seen as longer than usual response times and queued I/Os on the host. If the performance impact is too severe, the QoS limit may need to be adjusted.

Maximum limits are enforced only from I/O that arrives from an external host. These limits are not enforced on internal synchronous or asynchronous replication or migration I/O. QoS policies are also not replicated to a remote system. When replication is configured, it is suggested to configure the same QoS limits on each end of the replication session.

Having different limits in a metro volume configuration may change the pattern of I/O, as a host may prefer paths that have a higher QoS limit.

Dell Technologies recommends setting QoS on workloads which might monopolize system resources and starve other applications of their required performance. Consider some of the following opportunities to utilize PowerStore QoS:

- Limit the bandwidth available to large-block applications, such as backup jobs, which may be increasing the latency on other small-block workloads.
- Limit the IOPS capability of Thin Clones which are used in Test/Dev environments, so that they do not impact the capability of the associated source objects.
- Limit the IOPS / bandwidth that is available to non-essential applications that are hosted on the same system as your critical applications.

File storage resources

File storage resources are accessed through NAS protocols, such as NFS and SMB. A NAS server can provide access to a file system using all NAS protocols simultaneously, if configured for multiprotocol access.

Appliance balance for file workloads

A single NAS server uses compute resources from only one node of the appliance. It is recommended to create at least two NAS servers (one on each node) so that resources from both nodes contribute to file performance.

NAS servers can be manually moved from one node to the other. This action can be done to balance the workload if one node is busier than the other. All file systems that are served by a given NAS server move with the NAS server to the other node.

PowerStore features and layered applications

Introduction

This section discusses the features and layered applications available with PowerStoreOS.

Data reduction

PowerStore provides data-reduction capabilities such as zero-detect, compression, and deduplication. Data reduction is integrated into the PowerStore architecture and is always active. During periods of high write activity, PowerStore may defer the deduplication of data, and devote those resources to servicing the client workload. During periods of low activity, PowerStore will use excess resources to re-examine any data written during these periods for duplicates, to regain any space savings that were not initially realized.

Replication

PowerStore provides native synchronous and asynchronous replication solutions to protect data and help organizations meet business goals for both data availability and protection. Supported storage resources for native synchronous and asynchronous replication are volumes, volume groups, thin clones, NAS servers, and file systems. The replication itself uses iSCSI or the optimized Dell proprietary TCP-based replication protocol through Ethernet (LAN) connections that was introduced in PowerStoreOS 3.0.

When configuring replication, it is recommended that you use different physical ports for replication and front-end Ethernet storage traffic. A specific storage network for replication can be created and mapped to the desired physical ports for replication.

Metro and synchronous replication also require a low latency network (<5ms). As all host I/O must be written to the local and remote system before acknowledging the host, network speeds and the performance of the peer system can impact the latency encountered by the host. It is suggested that the local and remote system models be equal in case a failover of I/O must occur.

Metro Volume

In PowerStoreOS 3.0 and higher, PowerStore appliances support native metro volume replication. This provides synchronous replication of spanned block storage volumes exclusively for VMware VMFS Datastores in an active/active configuration across two PowerStore clusters in metro distance. In PowerStoreOS 4.0, volume group support has been added, along with Windows and Linux support. For more information about Metro Volume support, see the white paper [Dell PowerStore: Replication Technologies](#). For best practices for Metro Volume (latency, distance, uniform/non-uniform configurations, and so on), see the white paper [Dell PowerStore: Metro Volume](#). Both are available on the [PowerStore Info Hub](#).

Snapshots and thin clones

All storage resources in PowerStore are thinly provisioned and space efficient, including snapshots and thin clones. Creation of a snapshot or thin clone requires only a quick duplication of pointers. After this action, they behave as independent storage resources and do not impact the performance of the source resource.

Secure snapshots

In PowerStoreOS 3.5 and later, the optional secure snapshot setting allows you to create snapshots for volumes and volume groups. When the secure snapshot setting is enabled, the snapshot is protected from deletion until the retention period expires. This provides a cost-effective line of defense against ransom attacks and accidental deletion of snapshots, volumes, or volume groups. Note that after it is set, the retention period cannot be reduced by the user. It is therefore important to keep this in mind when defining it. Secure snapshots can only be terminated after they expire or by customer-authorized Dell Support.

Native PowerProtect DD Backup Integration

In PowerStoreOS 3.5 and later, users can back up volumes and volume groups directly to a PowerProtect DD series appliance running on premises or in the cloud. This feature eliminates the requirement of a backup host because all backup traffic is offloaded to the storage appliances. PowerStore uses the replication tagged storage ports to connect and transfer data to the PowerProtect DD appliance leveraging the DD Boost protocol. As with replication, it is recommended that you use different physical ports for data protection and front-end Ethernet storage traffic.

In PowerStoreOS 4.0, an enhancement to storage networks allows users to create multiple storage networks with defined purposes. These purposes include Storage (iSCSI), Storage (NVMe/TCP), and Replication. This feature allows storage networks for host connectivity, and dedicated networks for replication. When using PowerStore's native PowerProtect DD integration, it is required that the Replication purpose only be added to a single storage network within the cluster. The Storage (iSCSI) purpose must be

assigned to the network with the Replication purpose and be configured on a port/network that has connectivity to the remote PowerProtect appliance.

Cluster migrations

Block storage resources that are migrating between appliances in a cluster may see impacted performance from the migration activity. It is recommended to run migrations at a time when the resource is less busy.

PowerStoreOS upgrades

New versions of the PowerStoreOS are applied using a nondisruptive upgrade process. Because half of the system hardware resources are unavailable during parts of the upgrade, it is recommended to perform upgrades during planned maintenance windows. Alternately, perform upgrades when the system is less busy to minimize the impact to clients. Users should execute a Pre-Upgrade Health Check before their scheduled maintenance window to ensure that any issues can be resolved.

External host considerations

Introduction

This section highlights host configuration changes that may be necessary to access PowerStore volumes efficiently. For details about the appropriate settings for a host type, see the Host Connectivity Guides on [E-Lab Navigator](#).

Host configuration

Because PowerStore is a new class of storage, host operating systems may not natively recognize PowerStore volumes and apply the appropriate settings. For optimal performance, check that the appropriate configuration changes have been applied to all hosts that are connected to a PowerStore. The Host Connectivity Guides on [E-Lab Navigator](#) have recommendations for the following:

- MPIO settings: Path checker and timeout values
- iSCSI settings: Time-out and queue depth values; disabling delayed ACK
- Fibre Channel settings: Queue depth values
- Network settings: Jumbo frames and flow control
- Unmap operations
- VMware ESXi claim rules

For other recommended configurations for VMware ESXi and vSphere, see the document [Dell PowerStore: Virtualization Integration](#).

Host file systems

When a host is attached to a PowerStore block volume, the host can use this volume as a raw device, or it can create a local file system on the volume first. When a local file system is being created, it is recommended to disable SCSI unmap. When PowerStore creates a volume, all space is already unmapped; the host-based unmap is redundant and generates unnecessary load on PowerStore.

When creating a local file system, it is recommended to use a file system block size (allocation unit) that is an even multiple of 4 KB.

It is typically not necessary to perform alignment when creating a local file system. If alignment is performed, it is recommended to use an offset of 1 MB.

VMware

PowerStore is tightly integrated with VMware applications.

For other recommended configurations for VMware ESXi and vSphere, see the document [Dell PowerStore: Virtualization Integration](#).

Application considerations

PowerStore is well integrated with the most widely used enterprise applications. For best practice recommendations for specific applications, see the solutions-focused white papers available on the [PowerStore Info Hub](#).

Conclusion

Summary

This white paper provides configuration and usage recommendations for PowerStore in general use cases. For a detailed discussion of the reasoning or methodology behind these recommendations, or for additional guidance around more specific use cases, contact your Dell Technologies representative.

References

Dell Technologies documentation

The [Dell Technologies Info Hub > Storage](#) site provides expertise that helps to ensure customer success with Dell Technologies storage platforms.

[Dell.com/powerstoredocs](https://dell.com/powerstoredocs) provides detailed documentation about how to install, configure, and manage PowerStore systems.