Key Dates:

- Early quoting – May 8th
- DTW Keynote announcement - May 20th
- RTS May 21st

Key Sales Resources

Storage Central SharePoint – PowerScale Resources

Launch Overview and Goals

Q. What is the K2 launch?

A: K2 announces the launch of PowerScale F910, the newest addition to our next-generation all-flash line up that includes the F210 and F710.

These PowerScale nodes leverage Dell's newest PowerEdge architecture. They were designed and customized to deliver the highest performance for demanding workloads, such as generative AI, with the software-defined persistent memory, massive performance boosts and greater density.

Q. What is the RTS?

A: In the first wave, the F210 and F710 were launched in Feb.
In the second wave, the F910 will RTS on May 21st, 2024.

Q. When can I start selling the new F910 all-flash node?

A. Early quoting has opened on May 8th, with the RTS on May 21st.

Q. How do we refer to this node internally?

A: Internally we often refer to these nodes as "16G" because they are based on the "16G" PowerEdge Intel processor architecture. Externally we refer to them as next-generation PowerScale.

Similarly, you may have seen the reference to the current generation PowerScale nodes – F900, F600 and F200 as "14G" but we don't use these numbers externally

Q: What are the goals of this launch?

A. The goals are to:

- Position PowerScale as preferred storage platform for AI, as a key component of our AI-ready data platform
- Combat market pressure from competitors such as VAST, WEKA
- Maintain business continuity for existing customers.

Q. What is the market context?

A: Customers are looking for file storage that is highly performant, efficient and flexible so that it can support an entire range of workloads, from traditional file workloads such as file shares and file directories up to the most demanding AI/Gen AI and analytics workloads. The market trends include:

- Customers shifting to AI-first operating model and mindset, therefore need faster, secure, efficient and more economically viable infrastructure solutions
- Enterprise sustainability goals to reduce storage footprint with the best storage density possible

Next-generation nodes positioning and target workloads

Q: What is the high-level positioning for the nodes:

A: The new PowerScale nodes are positioned as:

- F210: an optimal platform for performance and small capacity requirements
- F710: a balanced platform for high performance and density in 1RU
- F910: optimized for capacity density in 2RU

Q: What makesthe F910 a key component of Dell's AI-ready data platform?

A: The PowerScale F910 stands as the newest cornerstone of our AI-ready data platform, delivering massive performance and density to scale AI workloads. Boasting up to 24 NVMe SSD drives per node and 300 PBs of effective storage per cluster, it ensures GPUs remain fully utilized for large-scale model training. Combined with the latest OneFS software, the F910 features federal-grade security, enabling customers to safeguard their data against threats like data poisoning and model inversion. Its multicloud flexibility lets users seamlessly manage AI workloads across various environments with consistent on-prem experience. Plus, it offers exceptional value to sustainably manage AI data growth while minimizing your storage footprint and power consumption.

Q: What workloads are these PowerScale all-flash nodes designed for?

A: These nodes are ideal for:

- Gen AI/AI in all phases of workflows
- Media & Entertainment (M&E) Content Creation & Rendering workflows
- Electronic Design Automation (EDA), for high concurrency & low latency workloads

As well as for the vertical industry workloads

- M&E
- EDA / Smart Factory / Automation
- Healthcare & Life Sciences
- Financial Services, High Frequency Trading
- Energy
- Education
- Transportation

Key features

Q. What are the general specs of the F910?

A: Here are the key specs:

- **F910 2RU**: requires OneFS 9.8; leverages PowerEdge R760; capacity in a minimum 6U cluster: min 276 TB – 2.2 PB raw

Q: What features of PowerEdge platform are being leveraged?

A: F910 leverage Dell PowerEdge R760 and include:

- 4th generation Intel Xeon Sapphire Rapids CPUs
- DDR5 – latest DRAM with greater speed and bandwidth
- PCIe Gen 5: 4X throughput compared to PCIe Gen 3 in the previous gen (14G)

Transition to next gen

Q. What is the transition from the current gen to the next gen nodes?

- F900 node customers can be transitioned to F910

Q. Will there be node-pool compatibility?

A: F910 is node-pool compatible with F900.

Q. What is the EOL plan for F900?

A: There will be a separate 411 EOL pre-announcement on this topic coming soon.

Key claims

Q. What are our key claims for these new nodes?

A: **Performance claims** based on the testing of F910 with OneFS 9.8 as compared with F900 on OneFS 9.5 (to provide a year over year gain) include:

- Up to 127% improved streaming write
- Up to 2x greater performance per watt
- Over 3 TB/s of read throughput per maximum cluster

Product claims for positioning F910 as the capacity node

- 20% greater node density per RU compared to the F710

For more information

Q. Where can I find the competitive positioning releted to this launch?

A. For the updated information on how we compete with our new all-flash nodes, please go to the UDS storage section on Klue.

Q: Where can I find more information?

A. Please check out the launch assets Storage Central SharePoint page:

- Prior to the RTS: NDA presentation, Ordering Guide and Sales FAQs
- At RTS: the complete set of launch internal and external assets.

Q: What is the sales and technical enablement plan for this launch?

A. Our first UDS Sales Enablement took place on 5/1. The recording and the presentation can be found here.